



# ‘Meme’ Stock Outlier Detection

MSDS 436 Final Project

Arthur Swanson  
Husein Adenwala



# Executive Summary

**Introduction:** Meme stocks have become very popular among retail investors that comment on social media platforms during the COVID-19 pandemic. Meme stocks are created when a company's shares start trending with individual investors on social media platforms and quickly skyrocket in price.

With a large amount of data being produced on social media platforms, it is difficult to identify the meme stock before or when the price takes off. Unsupervised algorithms aided by stock data can aid in identifying metrics of a meme stock.

**Results:** We provide evidence that these stocks display common stylized facts for the dynamics of price, trading volume, and social media activity. Understanding these properties helps investment analysts in their decisions.

**Conclusion:** Stock price was inversely correlated (mildly) to positive sentiment on Twitter. K-means clustering shows that Tesla was the most distinct cluster and could potentially suggest that we can detect meme stock activity using K-means clustering but more data over a longer period is needed.

# Research Objective



# Sourcing the Data

Data Source	Process
Twitter API	<p>Hourly batch scripts to get data using the Twitter development (free tier) API for a group of stocks limited by API count limitation.</p> <p>Stock symbols were used as key words to search for the tweets; however, this did capture many irrelevant tweets that were cleaned up during sentiment analysis.</p>
Yahoo Finance API	<p>Due to limits in number of requests on Y Finance API, we ran on-demand scripts on an EC2 instance to get Stock price attribution data.</p> <p>This step was done after sentiment analysis and data transformation as the Y Finance data is curated and did not require any pre-processing.</p>

# Data Pipeline Overview



## Twitter API

Using a Twitter developer account, Python API to get tweets tagged to stock symbol



## Y Finance API

Get min, max, average, volume, open, close, ticker, data timestamp during ETL process



## Database

Store the transformed data in AWS Aurora PostgreSQL



## Ingestion

Load tweet data in S3 bucket with batch scripts



## Prepare

Sentiment analysis on tweets.  
Merge Y Finance and Twitter data on stock symbol and date stamp

# Data Ingestion, Analysis and Preparation

- Hourly python batch scripts to collect Twitter data using API.
- Total of 83,000 tweets were collected with limited twitter development access.
- The tweets with date and timestamps were ingested in Amazon S3 in csv format.
- Used NLP pipeline in an Ec2 instance with python to filter irrelevant tweets, perform text preprocessing, tokenization and perform sentiment analysis using a Bert based transformer model.
- Used Yahoo Finance API to collect stock attribution data for the stocks and date and timestamps on post processed Twitter data.
- Yahoo Finance and Twitter sentiment data was loaded to AWS aurora (PostgreSQL).

# Data Governance

Only collected tweets on days where stock market was open for trading.

Collected tweets at a random sample and did not save user information.

# Methodology

## Overview of the System Design

- Twitter and Yahoo API for data

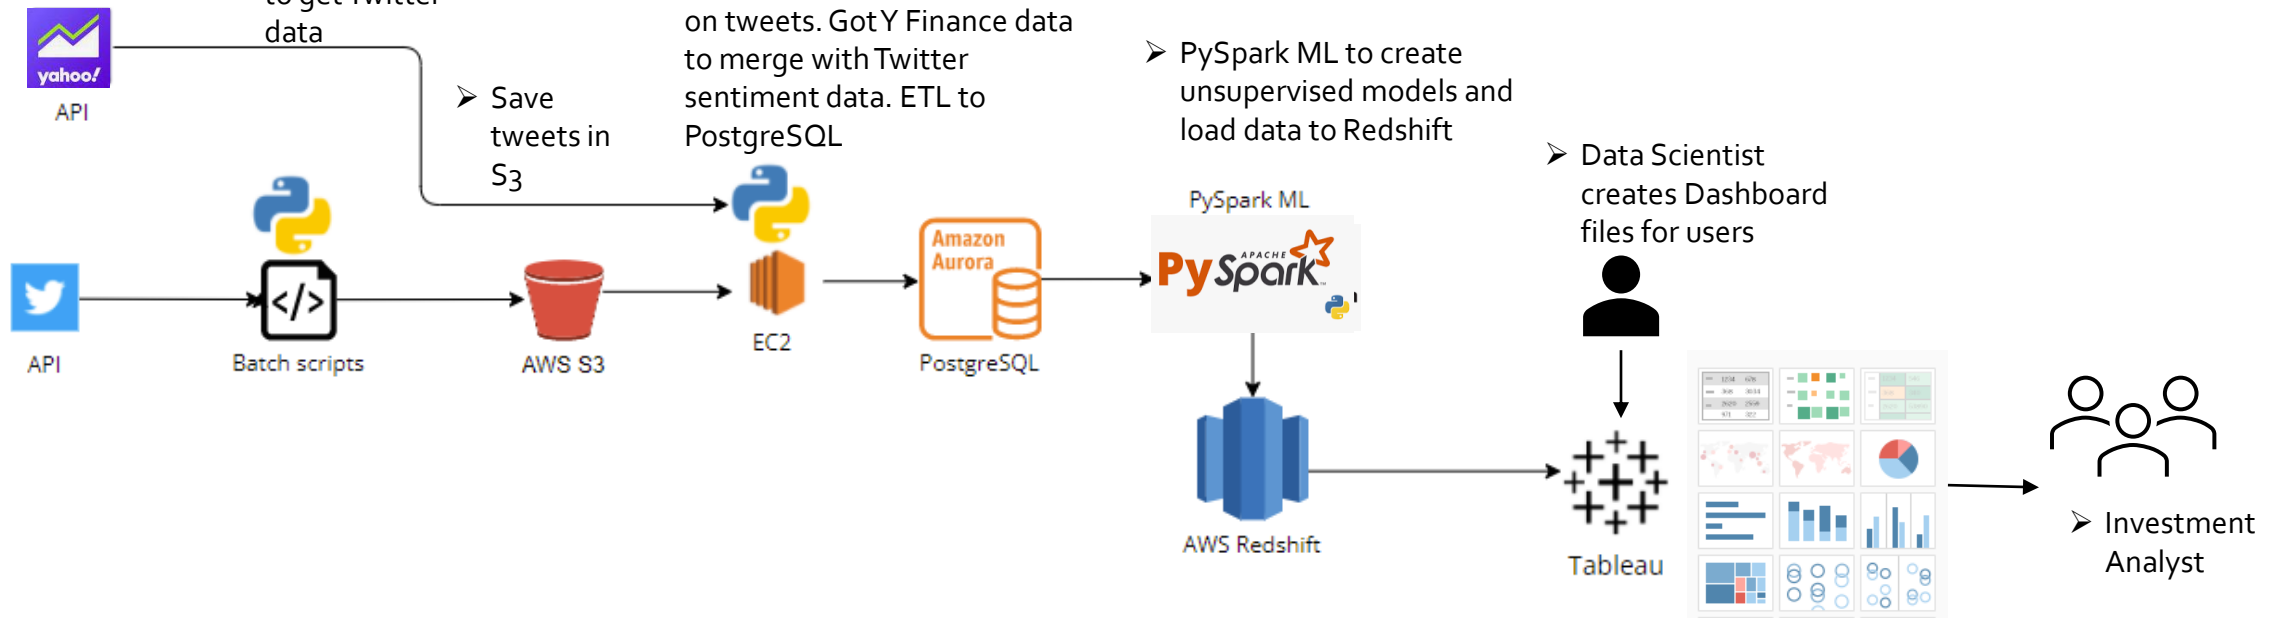
- Batch scripts on local environment to get Twitter data

- Scheduled Ec2 instance to perform sentiment analysis on tweets. GotY Finance data to merge with Twitter sentiment data. ETL to PostgreSQL

- PySpark ML to create unsupervised models and load data to Redshift

- Data Scientist creates Dashboard files for users

- Investment Analyst





# Methodology and Tools used

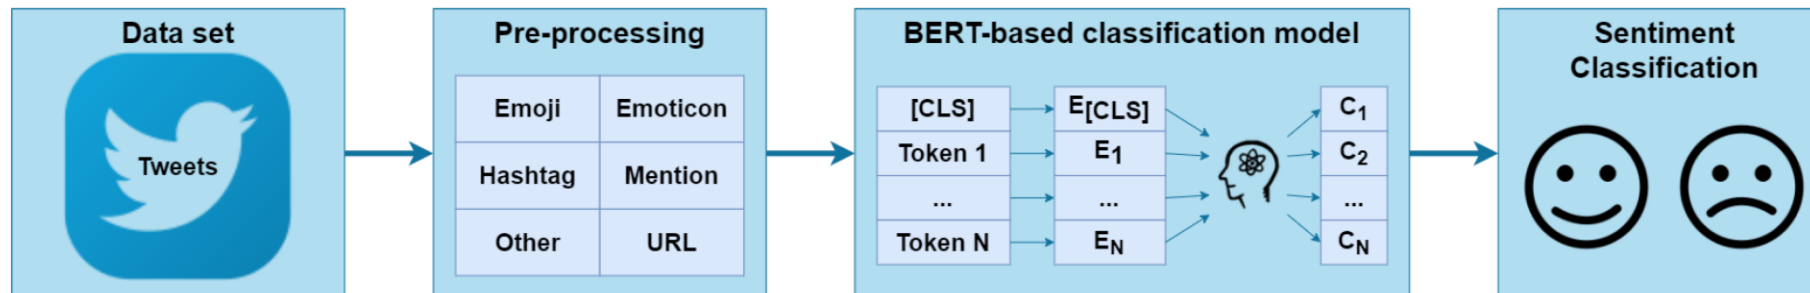
- Python to collect data from API
- Automation/Scheduling Scripts and Batch Processing
- Boto3 - Interface with Amazon APIs
- AWS S3 – Storage
- Ec2 instance with Cron scheduling python- ETL ( Data Transformation, NLP and getting Y Finance data)
- AWS Aurora - PostgreSQL to store the processed data
- PySpark ML – K-Means clustering model (DBSCAN did not provide discernable clusters)
- The transformed Yahoo and Twitter sentiment data with K-means clustering was pushed to AWS redshift
- AWS Redshift - Data Warehouse and Insight Result Storage for Deployment to Tableau
- Created Tableau TWB files that provide clustering results and dashboard for the end users

# NLP Sentiment Analysis

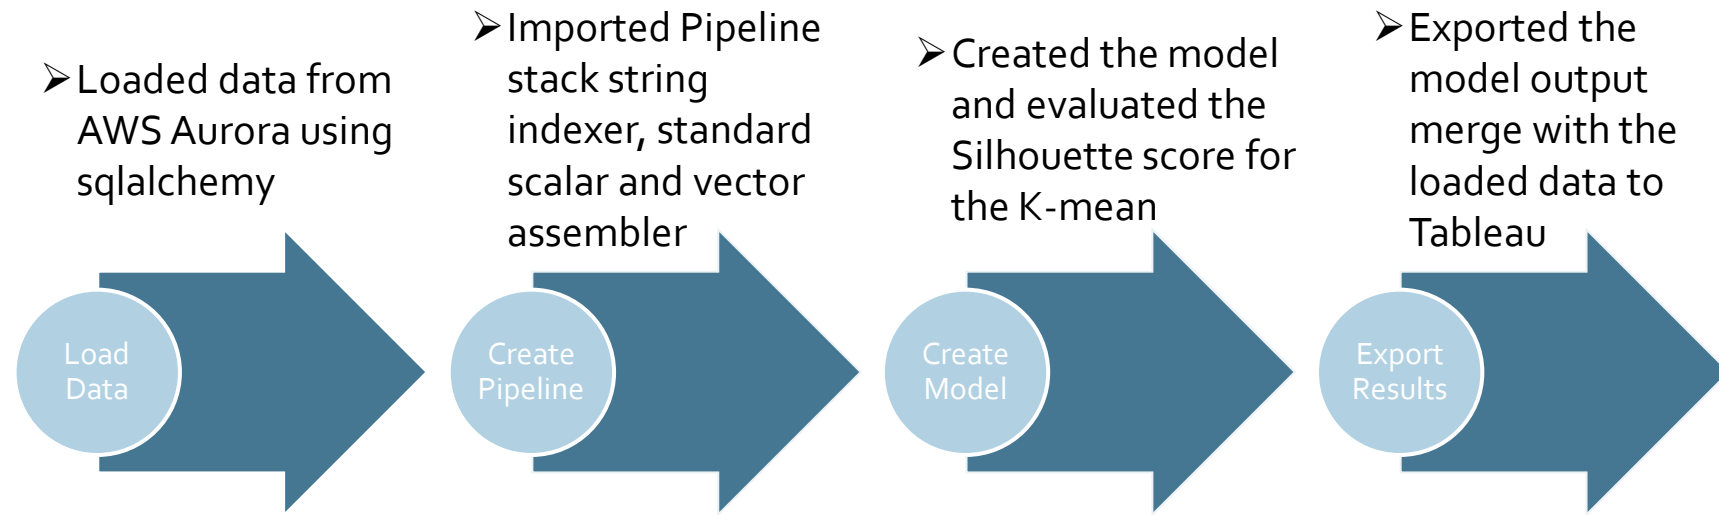
- Python NLTK library for cleaning text: Removed punctuation and special characters, lower cased, removed URLs and converted emoji to text.

- Used FLAIR NLP framework in python to tokenize and encode the text to get and predict the sentiment of the tweet (FLAIR uses a BERT based transformer model)

- Tweets were classified into positive and negative sentiments




# Unsupervised Models



# Results: Price – Sentiment Trend

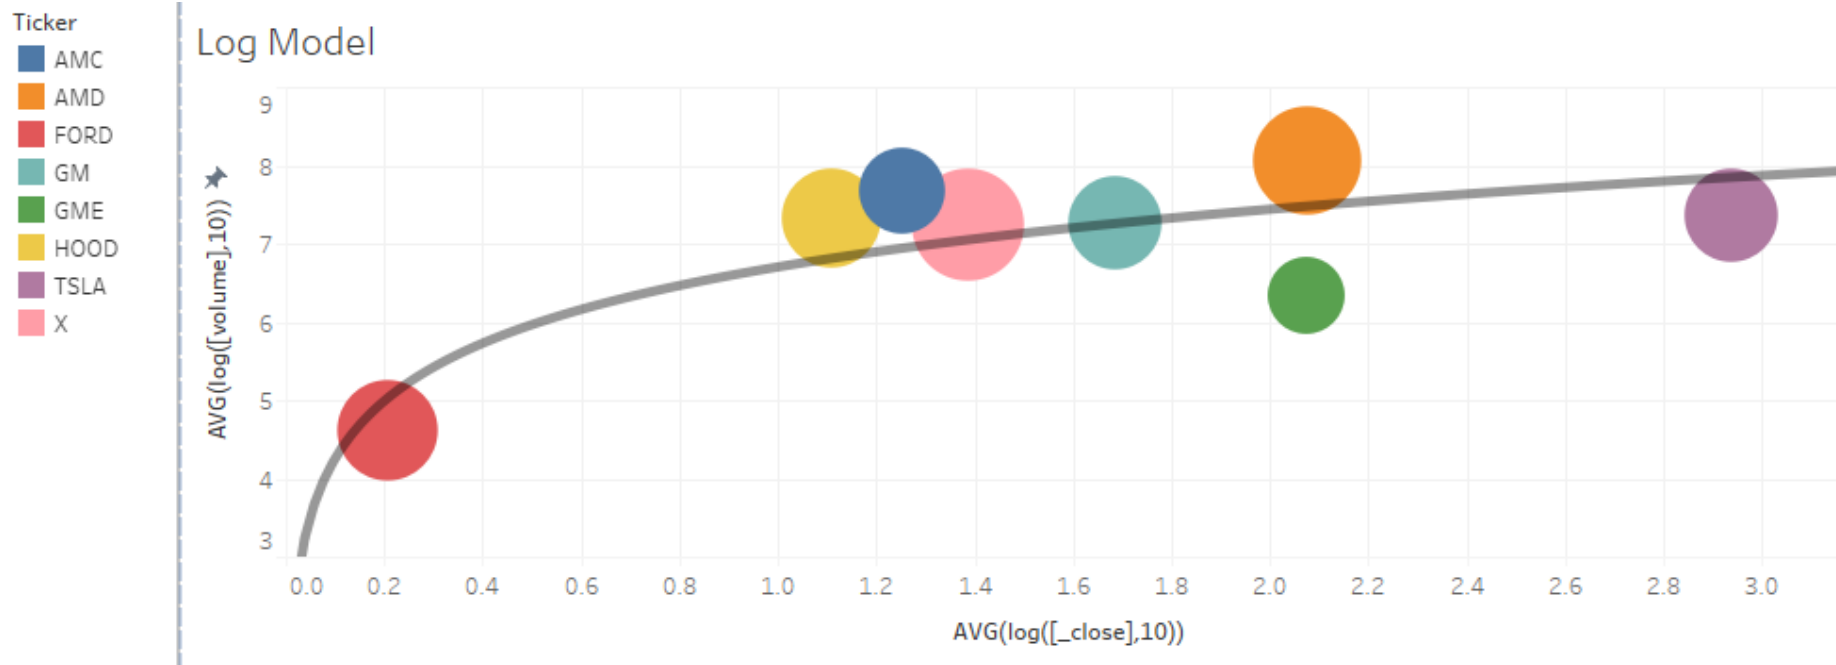
- Inverse (modest) relation between twitter sentiment and price
- Need more data over longer period to assess trend

Overall Trend



# Results: Price - Volume Trend


- Price and Volume on logarithmic trend.
- TSLA and FORD on extreme end.





# Results: K-Means Cluster

- TSLA can be identified as a distinct cluster
- However, FORD and TSLA were 'meme' stock and the cluster only detected TSLA



# Conclusion

## **Recommendations**

Using paid subscription, we can backlog tweets over longer period.

Get a better stock API for more resolute data over 5-minute windows.

Storing tweet and pricing in a NoSQL database like MongoDB would allow for flexibly in changing the schema as we gather more knowledge about the data requirements.

## **Lessons Learned**

Amazon EMR was not utilized and would have been a valuable feature to add.

Stream processing would be ideal application but could not be implemented.

Using all the processes in AWS would be useful for larger application.

# References

Das A (2021, Feb 10). K Means Clustering using PySpark on Big Data. Towards Data Science. Retrieved March 10 2022, from, <https://towardsdatascience.com/k-means-clustering-using-pyspark-on-big-data-6214beacdc8b>.

Jain S (2020, July 26). An Implementation of DBSCAN on PySpark. Towards Data Science. Retrieved March 10 2022 from, <https://towardsdatascience.com/an-efficient-implementation-of-dbscan-on-pyspark-3e2be646f57d>.

Magajna T (2018, Dec 24) Text Classification with State of the Art NLP Library — Flair. Towards Data Science. Retrieved March 10 20202, form <https://towardsdatascience.com/text-classification-with-state-of-the-art-nlp-library-flair-b541d7add21f>.

Galarnyk M (2017, Aug 2). T (2018, Dec 24) Text Classification with State of the Art NLP Library — Flair. Towards Data Science. Retrieved March 10 20202 form, <https://towardsdatascience.com/setting-up-and-using-jupyter-notebooks-on-aws-61a9648db6c5>.

# Appendix

# NLP Sentiment example

id	datetime	tweet	ticker	cleantweet	probability	sentiment
1490132285345218565	2022-02-06 01:16:36+00:00	RT @wolf_of_ape_st: \$AMC that historical OBV s...	AMC	RT of ape st AMC that historical OBV should te...	0.999317	NEGATIVE
1490132274532298752	2022-02-06 01:16:34+00:00	@Christalball93 AMC has nothing to do with DFV...	AMC	AMC has nothing to do with DFV you guys are to...	0.999816	NEGATIVE
1490132234975723520	2022-02-06 01:16:24+00:00	RT @EduardBrichuk: <i>AMC</i> at10k doesn't pose i...	AMC	RT AMC at 10k doesn t pose insane risk It ll s...	0.64562	POSITIVE
1490132217455996928	2022-02-06 01:16:20+00:00	\$AMC 🤔🤔 I believe trump 🇺🇸 <a href="https://t.co/H2qrB4lOyu">https://t.co/H2qrB4lOyu</a>	AMC	AMC I believe trump	0.802292	NEGATIVE
1490132212611788800	2022-02-06 01:16:19+00:00	RT @Christi92895977: #amc #AMCSQUEEZE #AMCSTOC...	AMC	RT amc AMCSQUEEZE AMCSTOCK	0.757214	NEGATIVE



# S3 Bucket example

## ► Account snapshot

Storage lens provides visibility into storage usage and activity trends. [Learn more](#)

## Buckets (9) [Info](#)

Buckets are containers for data stored in S3. [Learn more](#)



Copy ARN

Empty

Find buckets by name

	Name ▲	AWS Region ▼	Access ▼	Creation date
<input type="radio"/>	<a href="#">aws-assignment1-data</a>	US East (Ohio) us-east-2	<a href="#">Objects can be public</a>	January 26, 2022, 22:19:32 (UTC-06:00)
<input type="radio"/>	<a href="#">aws-assignment2-data</a>	US East (Ohio) us-east-2	<a href="#">Objects can be public</a>	January 29, 2022, 15:12:11 (UTC-06:00)
<input type="radio"/>	<a href="#">aws-big-data-ex3-test1</a>	US East (Ohio) us-east-2	<a href="#">Objects can be public</a>	January 27, 2022, 11:04:38 (UTC-06:00)
<input type="radio"/>	<a href="#">aws-big-data-ex3-test11</a>	US East (Ohio) us-east-2	<a href="#">Objects can be public</a>	January 27, 2022, 11:07:53 (UTC-06:00)
<input type="radio"/>	<a href="#">aws-big-data-ex3-test2</a>	US East (Ohio) us-east-2	<a href="#">Objects can be public</a>	January 27, 2022, 11:08:33 (UTC-06:00)
<input type="radio"/>	<a href="#">aws-emr-resources-180844430797-us-east-1</a>	US East (N. Virginia) us-east-1	<a href="#">Objects can be public</a>	February 13, 2022, 13:53:56 (UTC-06:00)
<input type="radio"/>	<a href="#">aws-logs-180844430797-us-east-1</a>	US East (N. Virginia) us-east-1	<a href="#">Objects can be public</a>	February 13, 2022, 13:49:28 (UTC-06:00)
<input type="radio"/>	<a href="#">msds436ex6</a>	US East (N. Virginia) us-east-1	<a href="#">Objects can be public</a>	February 13, 2022, 13:45:25 (UTC-06:00)
<input type="radio"/>	<a href="#">msdsgroup2data</a>	US East (Ohio) us-east-2	<a href="#">Objects can be public</a>	March 2, 2022, 10:11:03 (UTC-06:00)

# Amazon Aurora PostgreSQL instance

RDS > Databases

**Databases**

Filter by databases

	DB identifier	Role	Engine	Region & AZ	Size	Status
<input type="radio"/>	database-1	Regional cluster	Aurora PostgreSQL	us-east-2	1 instance	✓ Available
<input type="radio"/>	database-1-instance-1	Writer instance	Aurora PostgreSQL	us-east-2b	db.r6g.large	✓ Available

# Tableau Dashboard example

